


01 / HOW IT WORKS

The right project knowledge, for every model call.

The brain turns a project's reviewed documents, code, run records and approved memory into a bounded, hash-bound context bundle. The host decides what each role receives, checks it at the moment of the call, and keeps a receipt.



Who receives what


	TASK CONTRACT	BRAIN DOCS	CODE	RUN RECORDS	MEMORY
Planner	●	○	○	○	○
Implementer	●	○	○	○	○
Final auditor	●+	○	○	✗	○
Paired auditor	●+	○	○	✗	○

● always ○ only if the project grants it ✗ blocked in code + plus review evidence

DELIVERED TO EVERY CALL · 33 BUNDLES IN THE CAMPAIGN

13 first attempts, 7 retry deltas, 12 final audits, 1 paired audit. A final audit never gets a prior bundle, so it always starts fresh.

Retries resend only what changed




Real bundles. On retry the unchanged 11,251-byte code document becomes a 676-byte hash reference.

LOOKED UP ON DEMAND · `exact-lexical.v1`

singular context search | get: exact reference
+1000, exact phrase +100, +10 per matched token. It abstains below half the query tokens, and get refuses a stale version.

02 / MEMORY AND PROOF

Memory must be earned. Retrieval is measured.



Crash-safe by construction

```
.singular-memory/
├ records/      authoritative revision
├ operations/   journal: prepared → committed
├ checkpoints/ bounded recovery points
└ index.json    derived · rebuild any time
```

1. Lock, then finish any interrupted write.
2. Replaying an operation ID returns the stored answer.
3. Write the journal as *prepared*.
4. Apply only onto the expected predecessor.
5. Mark *committed*, then refresh the index.
6. Readers take no lock, write nothing, and refuse while a journal is pending.

The bug it closes. Replaying an interrupted approval once revived a tombstoned record.

Retrieval is measured (B5)

18 labeled cases · expected outcome

9 results	5 abstain	4 refuse
-----------	-----------	----------

abstain: says "not found" · refuse: stale, revoked, outside root, over budget

14/14

expected references included

0

incorrect selections

2

budget omissions, both declared

0

deviations from labels

Negative control: one wrong expectation turns the run to `deviated`, exit 4. No model call is needed. Not yet covered: memory sources.

Built node by node



Opt-in. Unconfigured projects are unchanged.

Documents can't grant themselves authority.

Project-local. No vector store or embeddings.

Readers never write. Safe in a sandbox.

Model output is a candidate, not truth.

Retirement sticks through every rebuild.

What is live today. Campaign runs used this path with code-only context: 33 admitted bundles. Brain documents, memory and lookups are proven in tests and the 18-case evaluation, not yet in a live campaign. That is the next milestone.